Homework for Session#7 : I was able to run the program on 8 nodes successfully using command: mpirun -n 8 python main.py

(2022-07-01//base) davidhe@thetagpu21:~/ai-science-training-series/ai-science-training-series/07_largeScaleTraining/src/ai4sci$ mpirun -n 8 python main.py


```
{
   "data": {
      "handler": "ilsvrc_dataset",
      "batch_size": 256,
      "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
      "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
      "shuffle_buffer": 200000,
      "reshuffle_each_iteration": true,
      "num_parallel_readers": 128,
      "prefetch_buffer_size": 8,
      "crop_image_size": [
         256,
         256
      ],
      "num_classes": 1000,
      "num_channels": 3
   }
}
{
   "data": {
      "handler": "ilsvrc_dataset",
      "batch_size": 256,
      "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
      "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
      "shuffle_buffer": 200000,
      "reshuffle_each_iteration": true,
      "num_parallel_readers": 128,
      "prefetch_buffer_size": 8,
      "crop_image_size": [
         256,
         256
```

```
      ],
      "num_classes": 1000,
      "num_channels": 3
    }
}
{
   "data": {
      "handler": "ilsvrc_dataset",
      "batch_size": 256,
      "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
      "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
      "shuffle_buffer": 200000,
      "reshuffle_each_iteration": true,
      "num_parallel_readers": 128,
      "prefetch_buffer_size": 8,
      "crop_image_size": [
         256,
         256
      ],
      "num_classes": 1000,
      "num_channels": 3
   }
}
{
   "data": {
      "handler": "ilsvrc_dataset",
      "batch_size": 256,
      "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
      "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
      "shuffle_buffer": 200000,
      "reshuffle_each_iteration": true,
      "num_parallel_readers": 128,
      "prefetch_buffer_size": 8,
      "crop_image_size": [
         256,
```

```
        256
      ],
      "num_classes": 1000,
      "num_channels": 3
    }
}
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 256,
    "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
    "crop_image_size": [
      256,
      256
    ],
    "num_classes": 1000,
    "num_channels": 3
  }
}
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 256,
    "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
    "crop_image_size": [
```

```
        256,
        256
      ],
      "num_classes": 1000,
      "num_channels": 3
    }
}
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 256,
    "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
    "crop_image_size": [
      256,
      256
    ],
    "num_classes": 1000,
    "num_channels": 3
  }
}
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 256,
    "train_filelist": "/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist": "/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
```

```
      "crop_image_size": [
        256,
        256
      ],
      "num_classes": 1000,
      "num_channels": 3
    }
}
```


num labels: 1000

build dataset /grand/projects/ALCFAITP/ilsvrc_train_filelist.txt

[2022-11-01 18:05:14,251][ai4sci.ilsvrc_dataset][INFO] - input filelist contains 1281167 files, estimated batches per rank 625

build dataset /grand/projects/ALCFAITP/ilsvrc_val_filelist.txt

[2022-11-01 18:05:16,720][ai4sci.ilsvrc_dataset][INFO] - input filelist contains 50000 files, estimated batches per rank 24

2022-11-01 18:05:17.397682: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397766: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397735: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397739: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397744: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397903: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397894: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

2022-11-01 18:05:17.397882: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

[2022-11-01 18:05:42,500][ai4sci.trainer][INFO] - [0] [0/1: 0 / 160145 (0%)] epoch=0.0000 dt=21.1748 running_loss=0.0000 batch_loss=0.0301 acc=0.0000 batch_acc=0.0000

[2022-11-01 18:05:48,889][ai4sci.trainer][INFO] - [0] [0/1: 2560 / 160145 (2%)] epoch=0.0000 dt=0.3209 running_loss=0.0005 batch_loss=0.0285 acc=0.0000 batch_acc=0.0000

[2022-11-01 18:05:55,277][ai4sci.trainer][INFO] - [0] [0/1: 5120 / 160145 (3%)] epoch=0.0000 dt=0.2636 running_loss=0.0010 batch_loss=0.0273 acc=0.0000 batch_acc=0.0000

[2022-11-01 18:06:02,591][ai4sci.trainer][INFO] - [0] [0/1: 7680 / 160145 (5%)] epoch=0.0000 dt=0.2570 running_loss=0.0014 batch_loss=0.0274 acc=0.0001 batch_acc=0.0000

[2022-11-01 18:06:11,307][ai4sci.trainer][INFO] - [0] [0/1: 10240 / 160145 (6%)] epoch=0.0000 dt=0.2835 running_loss=0.0019 batch_loss=0.0270 acc=0.0001 batch_acc=0.0000

[2022-11-01 18:06:18,505][ai4sci.trainer][INFO] - [0] [0/1: 12800 / 160145 (8%)] epoch=0.0000 dt=0.2875 running_loss=0.0023 batch_loss=0.0271 acc=0.0001 batch_acc=0.0000

[2022-11-01 18:06:23,924][ai4sci.trainer][INFO] - [0] [0/1: 15360 / 160145 (10%)] epoch=0.0000 dt=0.2402 running_loss=0.0027 batch_loss=0.0269 acc=0.0001 batch_acc=0.0000

[2022-11-01 18:06:31,435][ai4sci.trainer][INFO] - [0] [0/1: 17920 / 160145 (11%)] epoch=0.0000 dt=0.2606 running_loss=0.0032 batch_loss=0.0268 acc=0.0002 batch_acc=0.0078

[2022-11-01 18:06:39,450][ai4sci.trainer][INFO] - [0] [0/1: 20480 / 160145 (13%)] epoch=0.0000 dt=0.2753 running_loss=0.0036 batch_loss=0.0269 acc=0.0002 batch_acc=0.0039

[2022-11-01 18:06:50,207][ai4sci.trainer][INFO] - [0] [0/1: 23040 / 160145 (14%)] epoch=0.0000 dt=0.2650 running_loss=0.0040 batch_loss=0.0266 acc=0.0002 batch_acc=0.0039

[2022-11-01 18:06:55,590][ai4sci.trainer][INFO] - [0] [0/1: 25600 / 160145 (16%)] epoch=0.0000 dt=0.2569 running_loss=0.0044 batch_loss=0.0267 acc=0.0003 batch_acc=0.0039

[2022-11-01 18:07:02,452][ai4sci.trainer][INFO] - [0] [0/1: 28160 / 160145 (18%)] epoch=0.0000 dt=0.2776 running_loss=0.0049 batch_loss=0.0269 acc=0.0003 batch_acc=0.0039

[2022-11-01 18:07:09,539][ai4sci.trainer][INFO] - [0] [0/1: 30720 / 160145 (19%)] epoch=0.0000 dt=0.2394 running_loss=0.0053 batch_loss=0.0267 acc=0.0004 batch_acc=0.0039

[2022-11-01 18:07:15,832][ai4sci.trainer][INFO] - [0] [0/1: 33280 / 160145 (21%)] epoch=0.0000 dt=0.2726 running_loss=0.0057 batch_loss=0.0267 acc=0.0004 batch_acc=0.0000

[2022-11-01 18:07:23,048][ai4sci.trainer][INFO] - [0] [0/1: 35840 / 160145 (22%)] epoch=0.0000 dt=0.2579 running_loss=0.0061 batch_loss=0.0265 acc=0.0005 batch_acc=0.0000

[2022-11-01 18:07:30,525][ai4sci.trainer][INFO] - [0] [0/1: 38400 / 160145 (24%)] epoch=0.0000 dt=0.2607 running_loss=0.0066 batch_loss=0.0264 acc=0.0005 batch_acc=0.0000

[2022-11-01 18:07:39,912][ai4sci.trainer][INFO] - [0] [0/1: 40960 / 160145 (26%)] epoch=0.0000 dt=0.2507 running_loss=0.0070 batch_loss=0.0263 acc=0.0005 batch_acc=0.0000

[2022-11-01 18:07:48,019][ai4sci.trainer][INFO] - [0] [0/1: 43520 / 160145 (27%)] epoch=0.0000 dt=0.2615 running_loss=0.0074 batch_loss=0.0262 acc=0.0006 batch_acc=0.0078

[2022-11-01 18:07:55,563][ai4sci.trainer][INFO] - [0] [0/1: 46080 / 160145 (29%)] epoch=0.0000 dt=0.2537 running_loss=0.0078 batch_loss=0.0262 acc=0.0006 batch_acc=0.0039

[2022-11-01 18:08:03,165][ai4sci.trainer][INFO] - [0] [0/1: 48640 / 160145 (30%)] epoch=0.0000 dt=0.2542 running_loss=0.0082 batch_loss=0.0259 acc=0.0007 batch_acc=0.0078

[2022-11-01 18:08:10,243][ai4sci.trainer][INFO] - [0] [0/1: 51200 / 160145 (32%)] epoch=0.0000 dt=0.2505 running_loss=0.0087 batch_loss=0.0258 acc=0.0008 batch_acc=0.0078

[2022-11-01 18:08:15,860][ai4sci.trainer][INFO] - [0] [0/1: 53760 / 160145 (34%)] epoch=0.0000 dt=0.2521 running_loss=0.0091 batch_loss=0.0257 acc=0.0009 batch_acc=0.0000

[2022-11-01 18:08:23,063][ai4sci.trainer][INFO] - [0] [0/1: 56320 / 160145 (35%)] epoch=0.0000 dt=0.2761 running_loss=0.0095 batch_loss=0.0259 acc=0.0010 batch_acc=0.0000

[2022-11-01 18:08:31,042][ai4sci.trainer][INFO] - [0] [0/1: 58880 / 160145 (37%)] epoch=0.0000 dt=0.2520 running_loss=0.0099 batch_loss=0.0256 acc=0.0011 batch_acc=0.0117

[2022-11-01 18:08:37,980][ai4sci.trainer][INFO] - [0] [0/1: 61440 / 160145 (38%)] epoch=0.0000 dt=0.3157 running_loss=0.0103 batch_loss=0.0256 acc=0.0012 batch_acc=0.0156

[2022-11-01 18:08:43,720][ai4sci.trainer][INFO] - [0] [0/1: 64000 / 160145 (40%)] epoch=0.0000 dt=0.2561 running_loss=0.0107 batch_loss=0.0255 acc=0.0014 batch_acc=0.0078

[2022-11-01 18:08:50,812][ai4sci.trainer][INFO] - [0] [0/1: 66560 / 160145 (42%)] epoch=0.0000 dt=0.2689 running_loss=0.0111 batch_loss=0.0252 acc=0.0016 batch_acc=0.0273

[2022-11-01 18:09:03,561][ai4sci.trainer][INFO] - [0] [0/1: 69120 / 160145 (43%)] epoch=0.0000 dt=0.2667 running_loss=0.0115 batch_loss=0.0254 acc=0.0018 batch_acc=0.0078

[2022-11-01 18:09:08,951][ai4sci.trainer][INFO] - [0] [0/1: 71680 / 160145 (45%)] epoch=0.0000 dt=0.2694 running_loss=0.0119 batch_loss=0.0251 acc=0.0020 batch_acc=0.0078

[2022-11-01 18:09:16,254][ai4sci.trainer][INFO] - [0] [0/1: 74240 / 160145 (46%)] epoch=0.0000 dt=0.2568 running_loss=0.0123 batch_loss=0.0252 acc=0.0022 batch_acc=0.0156

[2022-11-01 18:09:24,450][ai4sci.trainer][INFO] - [0] [0/1: 76800 / 160145 (48%)] epoch=0.0000 dt=0.2667 running_loss=0.0127 batch_loss=0.0252 acc=0.0024 batch_acc=0.0078

[2022-11-01 18:09:31,591][ai4sci.trainer][INFO] - [0] [0/1: 79360 / 160145 (50%)] epoch=0.0000 dt=0.2654 running_loss=0.0131 batch_loss=0.0253 acc=0.0025 batch_acc=0.0078

[2022-11-01 18:09:38,735][ai4sci.trainer][INFO] - [0] [0/1: 81920 / 160145 (51%)] epoch=0.0000 dt=0.2363 running_loss=0.0135 batch_loss=0.0247 acc=0.0028 batch_acc=0.0039

[2022-11-01 18:09:44,431][ai4sci.trainer][INFO] - [0] [0/1: 84480 / 160145 (53%)] epoch=0.0000 dt=0.2488 running_loss=0.0139 batch_loss=0.0240 acc=0.0031 batch_acc=0.0469

[2022-11-01 18:09:52,108][ai4sci.trainer][INFO] - [0] [0/1: 87040 / 160145 (54%)] epoch=0.0000 dt=0.2676 running_loss=0.0143 batch_loss=0.0248 acc=0.0034 batch_acc=0.0156

[2022-11-01 18:09:59,868][ai4sci.trainer][INFO] - [0] [0/1: 89600 / 160145 (56%)] epoch=0.0000 dt=0.3179 running_loss=0.0147 batch_loss=0.0242 acc=0.0037 batch_acc=0.0078

[2022-11-01 18:10:07,014][ai4sci.trainer][INFO] - [0] [0/1: 92160 / 160145 (58%)] epoch=0.0000 dt=0.2626 running_loss=0.0150 batch_loss=0.0242 acc=0.0040 batch_acc=0.0273

[2022-11-01 18:10:12,308][ai4sci.trainer][INFO] - [0] [0/1: 94720 / 160145 (59%)] epoch=0.0000 dt=0.2854 running_loss=0.0154 batch_loss=0.0234 acc=0.0044 batch_acc=0.0195

[2022-11-01 18:10:19,587][ai4sci.trainer][INFO] - [0] [0/1: 97280 / 160145 (61%)] epoch=0.0000 dt=0.2671 running_loss=0.0158 batch_loss=0.0235 acc=0.0048 batch_acc=0.0234

[2022-11-01 18:10:27,318][ai4sci.trainer][INFO] - [0] [0/1: 99840 / 160145 (62%)] epoch=0.0000 dt=0.2922 running_loss=0.0162 batch_loss=0.0239 acc=0.0052 batch_acc=0.0078

[2022-11-01 18:10:33,983][ai4sci.trainer][INFO] - [0] [0/1: 102400 / 160145 (64%)] epoch=0.0000 dt=0.2408 running_loss=0.0166 batch_loss=0.0231 acc=0.0057 batch_acc=0.0469

[2022-11-01 18:10:43,231][ai4sci.trainer][INFO] - [0] [0/1: 104960 / 160145 (66%)] epoch=0.0000 dt=0.2853 running_loss=0.0169 batch_loss=0.0232 acc=0.0062 batch_acc=0.0273

[2022-11-01 18:10:49,501][ai4sci.trainer][INFO] - [0] [0/1: 107520 / 160145 (67%)] epoch=0.0000 dt=0.2350 running_loss=0.0173 batch_loss=0.0229 acc=0.0067 batch_acc=0.0352

[2022-11-01 18:10:59,129][ai4sci.trainer][INFO] - [0] [0/1: 110080 / 160145 (69%)] epoch=0.0000 dt=0.2678 running_loss=0.0177 batch_loss=0.0230 acc=0.0072 batch_acc=0.0273

[2022-11-01 18:11:04,687][ai4sci.trainer][INFO] - [0] [0/1: 112640 / 160145 (70%)] epoch=0.0000 dt=0.2541 running_loss=0.0180 batch_loss=0.0235 acc=0.0077 batch_acc=0.0234

[2022-11-01 18:11:12,459][ai4sci.trainer][INFO] - [0] [0/1: 115200 / 160145 (72%)] epoch=0.0000 dt=0.2701 running_loss=0.0184 batch_loss=0.0229 acc=0.0083 batch_acc=0.0273

[2022-11-01 18:11:20,629][ai4sci.trainer][INFO] - [0] [0/1: 117760 / 160145 (74%)] epoch=0.0000 dt=0.2857 running_loss=0.0188 batch_loss=0.0224 acc=0.0091 batch_acc=0.0391

[2022-11-01 18:11:28,233][ai4sci.trainer][INFO] - [0] [0/1: 120320 / 160145 (75%)] epoch=0.0000 dt=0.2670 running_loss=0.0191 batch_loss=0.0224 acc=0.0097 batch_acc=0.0352

[2022-11-01 18:11:35,077][ai4sci.trainer][INFO] - [0] [0/1: 122880 / 160145 (77%)] epoch=0.0000 dt=0.2695 running_loss=0.0195 batch_loss=0.0220 acc=0.0104 batch_acc=0.0430

[2022-11-01 18:11:40,518][ai4sci.trainer][INFO] - [0] [0/1: 125440 / 160145 (78%)] epoch=0.0000 dt=0.2442 running_loss=0.0198 batch_loss=0.0215 acc=0.0111 batch_acc=0.0469

[2022-11-01 18:11:47,833][ai4sci.trainer][INFO] - [0] [0/1: 128000 / 160145 (80%)] epoch=0.0000 dt=0.2619 running_loss=0.0202 batch_loss=0.0221 acc=0.0119 batch_acc=0.0352

[2022-11-01 18:11:55,759][ai4sci.trainer][INFO] - [0] [0/1: 130560 / 160145 (82%)] epoch=0.0000 dt=0.3036 running_loss=0.0205 batch_loss=0.0213 acc=0.0127 batch_acc=0.0703

[2022-11-01 18:12:03,315][ai4sci.trainer][INFO] - [0] [0/1: 133120 / 160145 (83%)] epoch=0.0000 dt=0.3040 running_loss=0.0209 batch_loss=0.0214 acc=0.0136 batch_acc=0.0820

[2022-11-01 18:12:11,366][ai4sci.trainer][INFO] - [0] [0/1: 135680 / 160145 (85%)] epoch=0.0000 dt=0.2705 running_loss=0.0212 batch_loss=0.0227 acc=0.0145 batch_acc=0.0469

[2022-11-01 18:12:16,800][ai4sci.trainer][INFO] - [0] [0/1: 138240 / 160145 (86%)] epoch=0.0000 dt=0.2551 running_loss=0.0216 batch_loss=0.0218 acc=0.0153 batch_acc=0.0469

[2022-11-01 18:12:23,904][ai4sci.trainer][INFO] - [0] [0/1: 140800 / 160145 (88%)] epoch=0.0000 dt=0.2677 running_loss=0.0219 batch_loss=0.0213 acc=0.0162 batch_acc=0.0742

[2022-11-01 18:12:31,633][ai4sci.trainer][INFO] - [0] [0/1: 143360 / 160145 (90%)] epoch=0.0000 dt=0.2629 running_loss=0.0223 batch_loss=0.0211 acc=0.0172 batch_acc=0.0273

[2022-11-01 18:12:42,813][ai4sci.trainer][INFO] - [0] [0/1: 145920 / 160145 (91%)] epoch=0.0000 dt=0.2506 running_loss=0.0226 batch_loss=0.0203 acc=0.0180 batch_acc=0.0664

[2022-11-01 18:12:48,292][ai4sci.trainer][INFO] - [0] [0/1: 148480 / 160145 (93%)] epoch=0.0000 dt=0.2529 running_loss=0.0229 batch_loss=0.0210 acc=0.0190 batch_acc=0.0664

[2022-11-01 18:12:56,383][ai4sci.trainer][INFO] - [0] [0/1: 151040 / 160145 (94%)] epoch=0.0000 dt=0.2713 running_loss=0.0233 batch_loss=0.0209 acc=0.0199 batch_acc=0.0664

[2022-11-01 18:13:03,683][ai4sci.trainer][INFO] - [0] [0/1: 153600 / 160145 (96%)] epoch=0.0000 dt=0.2718 running_loss=0.0236 batch_loss=0.0204 acc=0.0208 batch_acc=0.0625

[2022-11-01 18:13:13,029][ai4sci.trainer][INFO] - [0] [0/1: 156160 / 160145 (98%)] epoch=0.0000 dt=0.2563 running_loss=0.0239 batch_loss=0.0201 acc=0.0219 batch_acc=0.0898

[2022-11-01 18:13:20,100][ai4sci.trainer][INFO] - [0] [0/1: 158720 / 160145 (99%)] epoch=0.0000 dt=0.2627 running_loss=0.0243 batch_loss=0.0200 acc=0.0231 batch_acc=0.0508

[2022-11-01 18:13:23,729][__main__][INFO] - Total training time: 492.14545607566833 seconds

[2022-11-01 18:13:23,729][__main__][INFO] - Total training time: 492.14706540107727 seconds

[2022-11-01 18:13:23,729][__main__][INFO] - Total training time: 492.1463449001312 seconds

[2022-11-01 18:13:23,730][__main__][INFO] - Average time per epoch in the last 5: 486.8954539299011

[2022-11-01 18:13:23,730][__main__][INFO] - Total training time: 492.1457154750824 seconds

[2022-11-01 18:13:23,730][__main__][INFO] - Total training time: 492.14664220809937 seconds

[2022-11-01 18:13:23,730][__main__][INFO] - Total training time: 492.1481206417084 seconds

[2022-11-01 18:13:23,730][__main__][INFO] - Total training time: 492.1473476886749 seconds

[2022-11-01 18:13:23,730][__main__][INFO] - Average time per epoch in the last 5: 486.8939199447632

[2022-11-01 18:13:23,730][__main__][INFO] - Average time per epoch in the last 5: 486.89657974243164

[2022-11-01 18:13:23,730][__main__][INFO] - Average time per epoch in the last 5: 486.89453506469727

[2022-11-01 18:13:23,730][__main__][INFO] - Average time per epoch in the last 5: 486.89521288871765

[2022-11-01 18:13:23,730][__main__][INFO] - Average time per epoch in the last 5: 486.8961760997772

[2022-11-01 18:13:23,731][__main__][INFO] - Average time per epoch in the last 5: 486.89490699768066

2022-11-01 18:13:23.768105: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr { key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6 } dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape { dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" model: "241" frequency: 2245 num_cores: 32 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { key: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 } outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }

[2022-11-01 18:13:36,622][__main__][INFO] - --------------------

[2022-11-01 18:13:36,622][__main__][INFO] - [TEST] Accuracy: 0%

[2022-11-01 18:13:36,622][__main__][INFO] - --------------------

[2022-11-01 18:13:36,625][__main__][INFO] - ----------------------------

[2022-11-01 18:13:36,626][__main__][INFO] - [TRAIN] loss=0.0031 acc=0%

[2022-11-01 18:13:36,626][__main__][INFO] - ----------------------------

[2022-11-01 18:13:37,173][__main__][INFO] - Total training time: 505.589248418808 seconds

[2022-11-01 18:13:37,174][__main__][INFO] - Average time per epoch in the last 5: 486.8990716934204

(2022-07-01//base) davidhe@thetagpu21:~/ai-science-training-series/ai-science-training-series/07_largeScaleTraining/src/ai4sci$